

Varit Kobutra

Professor Patricia McManus

Computer Vision ITAI 1378: 12461

20 June 2024

Reflective Journal: L05 Image Classification with SVM

This lab introduced me to the fundamentals of image classification using Support Vector Machines (SVM) and the CIFAR-10 dataset. It deepened my understanding of machine learning principles and illuminated the practical challenges inherent in model training. Key takeaways include:

Understanding SVM and Image Classification

- The SVM algorithm is a robust tool for identifying optimal decision boundaries between different classes of data (Burges 121).
- In image classification, SVM's strength lies in its ability to transform image data into a higher-dimensional space, revealing intricate patterns and relationships that are otherwise difficult to discern (Cortes and Vapnik 273).

Data Preparation, Model Training, and Evaluation

- The CIFAR-10 dataset was loaded, preprocessed (visualization, grayscale conversion, flattening), and split into training and testing sets.
- Challenges were encountered with long training times, both locally and on Google Colab, preventing full training as per the provided notebook instructions.
- To address this, I explored alternative approaches on Kaggle and GitHub, gaining valuable insights into hyperparameter optimization and model training strategies.
- Model performance was assessed using accuracy on the testing set, and predictions were visualized to understand strengths and weaknesses.

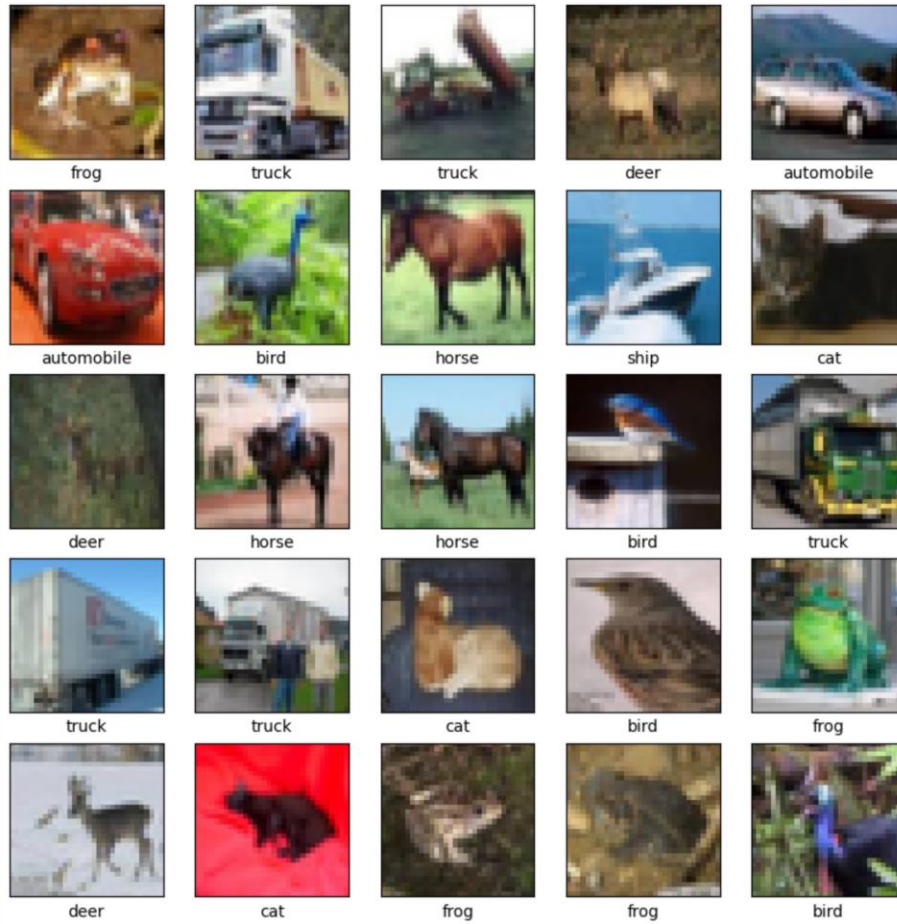
Challenges and Insights

- Hyperparameter tuning was explored to optimize accuracy, emphasizing the importance of understanding their impact on model performance.
- Principal Component Analysis (PCA) was applied to reduce dimensionality and accelerate training, highlighting the trade-off between computational efficiency and model accuracy (resulting accuracy: 31.39%).

Proof of Work

```
# Visualize some images
plt.figure(figsize=(10, 10))
for i in range(25):
    plt.subplot(5, 5, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i])
    plt.xlabel(class_names[y_train[i][0]])
plt.show()
```

✓ 0.2s



```
# Standardize the data
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train_split)
x_val_scaled = scaler.transform(x_val_split)
x_test_scaled = scaler.transform(x_test_flat)

# Apply PCA for dimensionality reduction
pca = PCA(n_components=100) # Reduce to 100 principal components
x_train_pca = pca.fit_transform(x_train_scaled)
x_val_pca = pca.transform(x_val_scaled)
x_test_pca = pca.transform(x_test_scaled)
```

✓ 14.2s

```
# Train the SVM classifier
svm_classifier = SVC(kernel='linear') # Linear kernel for faster training
svm_classifier.fit(x_train_pca, y_train_split.ravel())
```

✓ 25m 36.0s

SVC

SVC(kernel='linear')

```
# Make predictions on the testing set
y_pred = svm_classifier.predict(x_test_pca)

# Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
print("Classification Report:\n", classification_report(y_test, y_pred))
```

✓ 14.3s

Accuracy: 31.21%

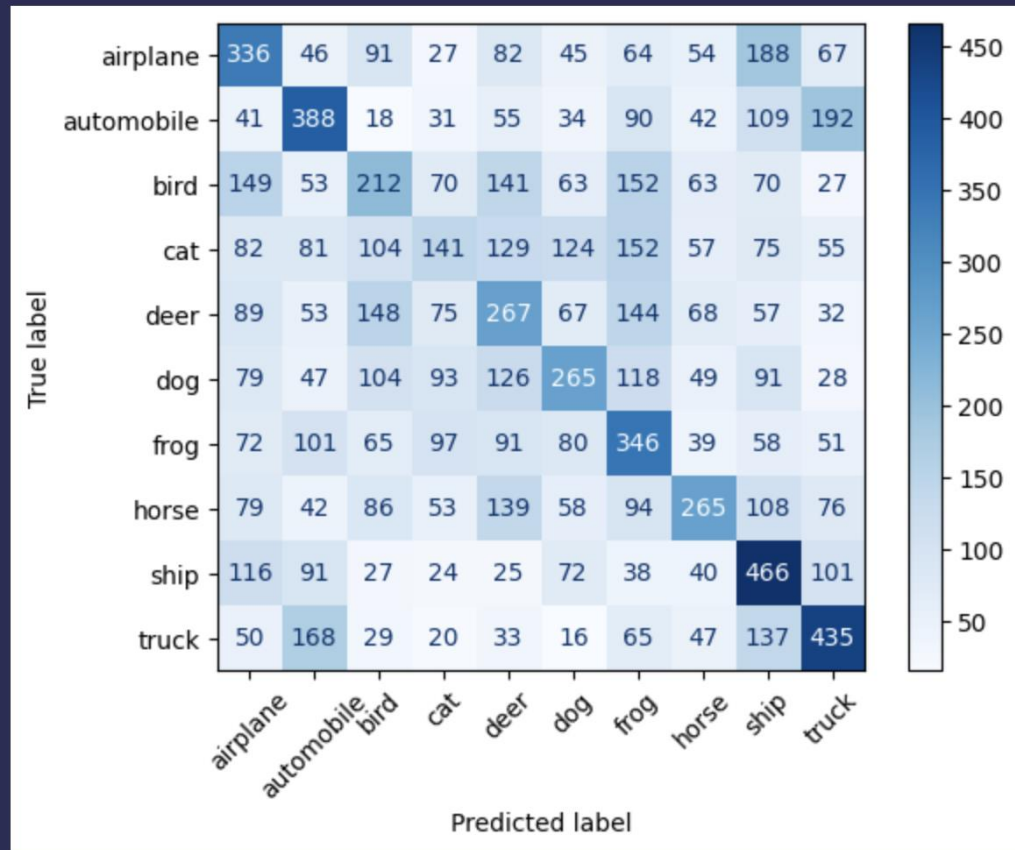
Classification Report:

	precision	recall	f1-score	support
0	0.31	0.34	0.32	1000
1	0.36	0.39	0.37	1000
2	0.24	0.21	0.23	1000
3	0.22	0.14	0.17	1000
4	0.25	0.27	0.26	1000
5	0.32	0.27	0.29	1000
6	0.27	0.35	0.31	1000
7	0.37	0.27	0.31	1000
8	0.34	0.47	0.40	1000
9	0.41	0.43	0.42	1000
accuracy			0.31	10000
macro avg	0.31	0.31	0.31	10000
weighted avg	0.31	0.31	0.31	10000

```
# Compute the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
disp.plot(cmap=plt.cm.Blues)
plt.xticks(rotation=45) # Rotate the x labels for better readability
plt.show()
```

✓ 0.1s



Works Cited

Burges, Christopher J.C. "A Tutorial on Support Vector Machines for Pattern Recognition." Data Mining and Knowledge Discovery, vol. 2, no. 2, 1998, pp. 121-167.

Cortes, Corinna, and Vladimir Vapnik. "Support-Vector Networks." Machine Learning, vol. 20, no. 3, 1995, pp. 273-297.

Krizhevsky, Alex. Learning Multiple Layers of Features from Tiny Images. Tech report, University of Toronto, 2009. <https://www.cs.toronto.edu/~kriz/cifar.html>